

▼ Task 1. Data Summarization:

Calculate basic summary statistics (mean, median, standard deviation, etc.) for each numerical variable.

```
from google.colab import files
```

```
uploaded = files.upload()
```

Choose Files healthcare-...ke-data.csv

- **healthcare-dataset-stroke-data.csv**(text/csv) - 316971 bytes, last modified: 1/26/2021 - 100% done
Saving healthcare-dataset-stroke-data.csv to healthcare-dataset-stroke-data.csv

```
import pandas as pd
```

```
dataset = pd.read_csv('healthcare-dataset-stroke-data.csv')
dataset.head()
```



	id	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_glucose_level	bmi	smoking_status
0	9046	Male	67.0	0	1	Yes	Private	Urban	228.69	36.6	formerly smoked
1	51676	Female	61.0	0	0	Yes	Self-employed	Rural	202.21	NaN	never smoked
2	31112	Male	80.0	0	1	Yes	Private	Rural	105.92	32.5	never smoked
3	60182	Female	49.0	0	0	Yes	Private	Urban	171.23	34.4	smoker
4	1665	Female	79.0	1	0	Yes	Self-employed	Rural	174.12	24.0	never smoked

```
numerical_variable = [feature for feature in dataset.columns if dataset[feature].dtypes != 'O']
print('Number of numerical variables: ', len(numerical_variable))
dataset[numerical_variable].head()
```

Number of numerical variables: 7

	id	age	hypertension	heart_disease	avg_glucose_level	bmi	stroke
0	9046	67.0	0	1	228.69	36.6	1
1	51676	61.0	0	0	202.21	NaN	1
2	31112	80.0	0	1	105.92	32.5	1
3	60182	49.0	0	0	171.23	34.4	1
4	1665	79.0	1	0	174.12	24.0	1

```
means = dataset.mean()
means
```

```
<ipython-input-4-2e51b005e08c>:1: FutureWarning: The default value of numeric_only in DataFrame.mean is deprecated. In a future version, it will default to False. In addition, specifying 'num
means = dataset.mean()
```

```

id          36517.829354
age         43.226614
hypertension 0.097456
heart_disease 0.054012
avg_glucose_level 106.147677
bmi         28.893237
stroke      0.048728
dtype: float64

```

Conclusion:

Based on the calculated means of features from the dataset:

ID: The average ID value is approximately 36,517.83, suggesting that the dataset may have a wide range of unique identifiers.

Age: The average age is approximately 43.23 years, indicating that the dataset contains individuals with a broad age range.

Hypertension: The mean hypertension rate is around 9.75%, implying that a relatively small proportion of individuals in the dataset have hypertension.

Heart Disease: The mean heart disease rate is approximately 5.40%, indicating that a small percentage of individuals in the dataset have a history of heart disease.

Average Glucose Level: The average glucose level is about 106.15 mg/dL, providing insights into the typical blood glucose levels in the dataset.

BMI (Body Mass Index): The mean BMI is roughly 28.89, representing the average body mass index of individuals in the dataset.

Stroke: The mean stroke rate is approximately 4.87%, indicating that a small proportion of individuals in the dataset have experienced a stroke.

These mean values provide a summary of the central tendencies of the dataset's features.

```

medians = dataset.median()
medians

```

```

<ipython-input-5-fbf1ff32e499>:1: FutureWarning: The default value of numeric_only in DataFrame.median is deprecated. In a future version, it will default to False. In addition, specifying 'n
medians = dataset.median()
id          36932.000
age         45.000
hypertension 0.000
heart_disease 0.000
avg_glucose_level 91.885
bmi         28.100
stroke      0.000
dtype: float64

```

For id, it represents the median identification number, which doesn't provide meaningful statistical information.

age has a median of 45.000, indicating that approximately half of the individuals in the dataset are below 45 years of age, and the other half are above 45.

hypertension, heart_disease, and stroke are binary variables, and their medians of 0.000 suggest that the majority of individuals do not have hypertension, heart disease, or stroke based on this dataset.

avg_glucose_level has a median of 91.885, which represents the middle value of the glucose levels in the dataset. This can be used to understand the typical glucose level of the population.

bmi has a median of 28.100, indicating that the middle value of BMI (Body Mass Index) in the dataset is 28.100, which can help assess the average weight status of individuals in the dataset.

```
standard_deviation = dataset.std()
standard_deviation
```

```
<ipython-input-6-45780231bd12>:1: FutureWarning: The default value of numeric_only in DataFrame.std is deprecated. In a future version, it will default to False. In addition, specifying 'nume
standard_deviation = dataset.std()
id                21161.721625
age               22.612647
hypertension      0.296607
heart_disease     0.226063
avg_glucose_level 45.283560
bmi               7.854067
stroke           0.215320
dtype: float64
```

age: The standard deviation of age is relatively low (22.61), indicating that the ages in the dataset are relatively close to the mean, with limited variation.

hypertension and heart_disease: Both hypertension and heart disease have low standard deviations (0.30 and 0.23, respectively), suggesting that these binary variables have little variability within the dataset.

avg_glucose_level: The standard deviation for average glucose level is relatively high (45.28), indicating a wider range of values, suggesting more significant variability.

bmi: The standard deviation for BMI (body mass index) is moderately high (7.85), indicating a notable spread in BMI values across the dataset.

stroke: The standard deviation for the target variable, 'stroke,' is relatively low (0.22), indicating that it is imbalanced, with a majority of non-stroke cases.

```
dataset.describe()
```

	id	age	hypertension	heart_disease	avg_glucose_level	bmi	stroke
count	5110.000000	5110.000000	5110.000000	5110.000000	5110.000000	4909.000000	5110.000000
mean	36517.829354	43.226614	0.097456	0.054012	106.147677	28.893237	0.048728
std	21161.721625	22.612647	0.296607	0.226063	45.283560	7.854067	0.215320
min	67.000000	0.080000	0.000000	0.000000	55.120000	10.300000	0.000000
25%	17741.250000	25.000000	0.000000	0.000000	77.245000	23.500000	0.000000
50%	36932.000000	45.000000	0.000000	0.000000	91.885000	28.100000	0.000000
75%	54682.000000	61.000000	0.000000	0.000000	114.090000	33.100000	0.000000
max	72940.000000	82.000000	1.000000	1.000000	271.740000	97.600000	1.000000

Count the frequency of unique values for categorical variables

```
categorical_variable = [feature for feature in dataset.columns if dataset[feature].dtypes == 'O']
print('Number of Categorical variables: ', len(categorical_variable))
dataset[categorical_variable].head()
```

Number of Categorical variables: 5

	gender	ever_married	work_type	Residence_type	smoking_status
0	Male	Yes	Private	Urban	formerly smoked
1	Female	Yes	Self-employed	Rural	never smoked
2	Male	Yes	Private	Rural	never smoked
3	Female	Yes	Private	Urban	smokes
4	Female	Yes	Self-employed	Rural	never smoked

```
for feature in categorical_variable:
    unique_value = dataset[feature].value_counts()
    print(f"Unique Values for {feature}:\n{unique_value}\n")
```

Unique Values for gender:

```
Female    2994
Male      2115
Other         1
Name: gender, dtype: int64
```

Unique Values for ever_married:

```
Yes    3353
No     1757
Name: ever_married, dtype: int64
```

Unique Values for work_type:

```
Private    2925
Self-employed    819
children    687
Govt_job     657
Never_worked    22
Name: work_type, dtype: int64
```

Unique Values for Residence_type:

```
Urban    2596
Rural    2514
Name: Residence_type, dtype: int64
```

Unique Values for smoking_status:

```
never smoked    1892
Unknown         1544
formerly smoked    885
smokes           789
Name: smoking_status, dtype: int64
```

Calculate the number of missing values for each variable

```
missing_values = dataset.isna().sum()
print(missing_values)
```




```
id            0
gender        0
age           0
hypertension  0
heart_disease 0
ever_married  0
work_type     0
Residence_type 0
avg_glucose_level 0
bmi           201
smoking_status 0
stroke        0
dtype: int64
```




▼ Task 2. Data Visualization:

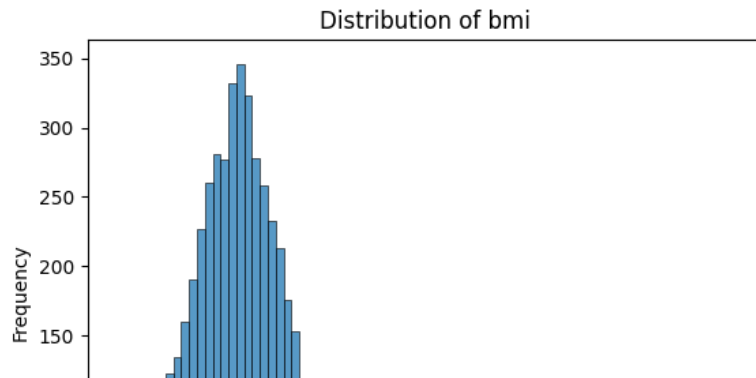

Create histograms or density plots to visualize the distribution of numerical variables

```
import seaborn as sns
import matplotlib.pyplot as plt

for Feature in numerical_variable:
    sns.histplot(dataset[Feature])
    plt.title(f'Distribution of {Feature}')
    plt.xlabel(Feature)
    plt.ylabel('Frequency')
    plt.show()
```














Generate bar plots or pie charts to visualize the distribution of categorical variables

```
for Feature in categorical_variable:
    unique_values = dataset[Feature].value_counts()
    sns.barplot(x=unique_values.index, y=unique_values.values)
    plt.title(f'Distribution of {Feature}')
    plt.xlabel(Feature)
    plt.ylabel('Frequency')
    plt.xticks(rotation=90)
    plt.show()
```







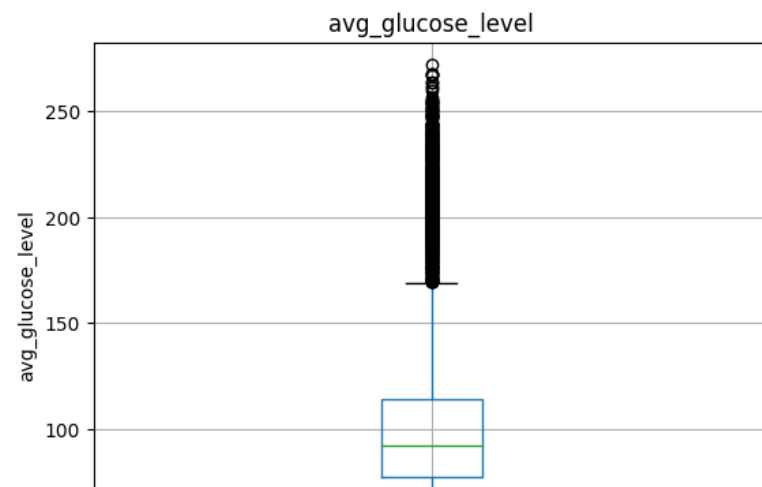







Create box plots to identify outliers and understand the spread of data

```
import numpy as np
for feature in numerical_variable:
    dataset.boxplot(column=feature)
    plt.ylabel(feature)
    plt.title(feature)
    plt.show()
```









Construct scatter plots to explore relationships between pairs of variables.

```
sns.pairplot(dataset[numerical_variable])  
plt.show()
```



Use heatmaps to visualize correlations between variables.

```
correlation_matrix = dataset[numerical_variable].corr()
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Correlation Heatmap')
plt.show()
```

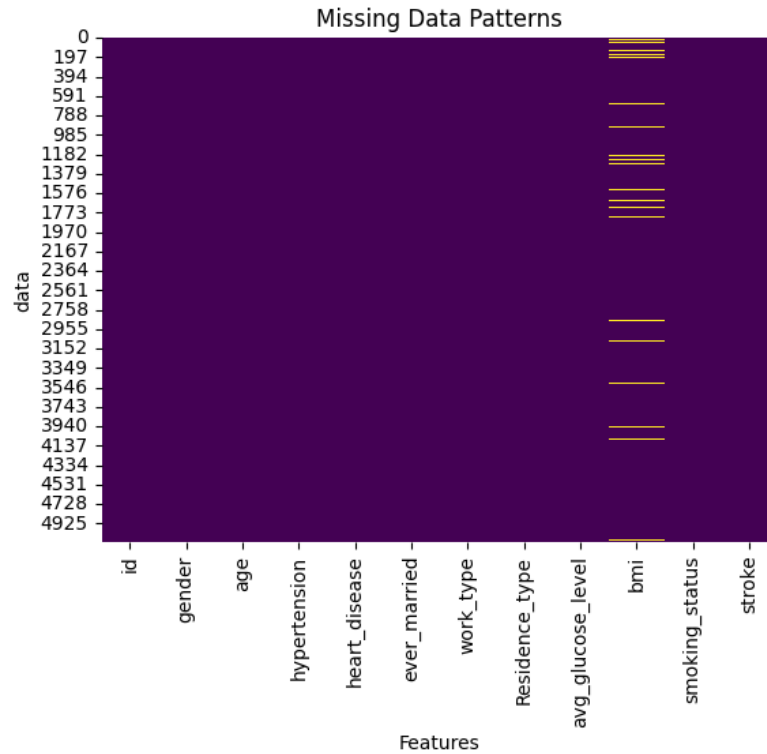


▼ Task 3. Handling Missing Data:

Explore the patterns of missing data across variables


```
missing_data_matrix = dataset.isnull()

sns.heatmap(missing_data_matrix, cmap='viridis', cbar=False)
plt.title('Missing Data Patterns')
plt.xlabel('Features')
plt.ylabel('data')
plt.show()
```



Decide on an appropriate strategy for handling missing values (imputation, removal, etc.)



We can fill the missing values by calculating distances with other samples and fill the missing value with closest neighbour.

▼ Task 4. Outlier Detection and Treatment:

Identify and visualize outliers in numerical variables.

```
for Feature in numerical_variable:
    data=dataset.copy()
```

```
if 0 in data[Feature].unique():  
    pass  
else:  
    data[Feature]=np.log(data[Feature])  
    data.boxplot(column=Feature)  
    plt.ylabel(Feature)  
    plt.title(Feature)  
    plt.show()
```



avg_glucose_level

Decide whether to remove, transform, or treat outliers based on domain knowledge and analysis goals.



Since data is belongs to medical domain and we cannot simply discard missing value instead we can normalize by taking log function



▼ Task 5. Data Distribution Analysis:


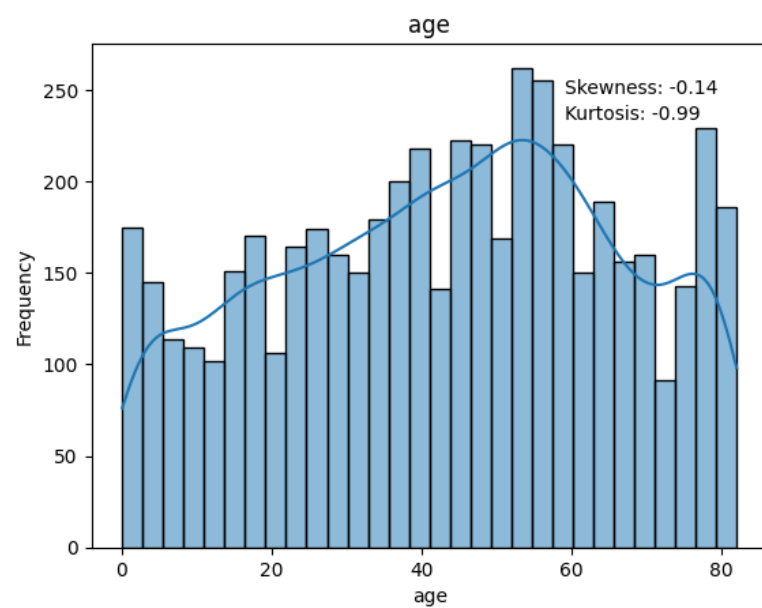




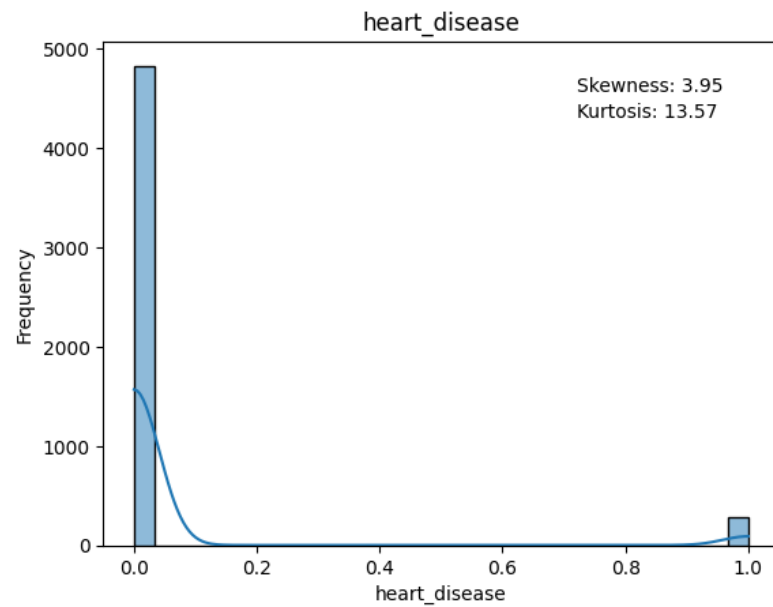

Visualize the data distribution and assess skewness and kurtosis




from scipy.stats import skew, kurtosis

```
for Feature in numerical_variable:
    sns.histplot(dataset[Feature], kde=True, bins=30)
    skewness = skew(dataset[Feature])
    kurt = kurtosis(dataset[Feature])
    plt.text(0.7, 0.9, f'Skewness: {skewness:.2f}', transform=plt.gca().transAxes)
    plt.text(0.7, 0.85, f'Kurtosis: {kurt:.2f}', transform=plt.gca().transAxes)
    plt.title(f'{Feature}')
    plt.xlabel(Feature)
    plt.ylabel('Frequency')
    plt.show()
```















▼ Task 6. Bivariate Analysis:




Analyze relationships between pairs of variables through scatter plots




```
for i in range(len(numerical_variable)):
    for j in range(i + 1, len(numerical_variable)):
        sns.scatterplot(data=dataset, x=numerical_variable[i], y=numerical_variable[j])
        plt.xlabel(numerical_variable[i])
        plt.ylabel(numerical_variable[j])
        plt.title(f'{numerical_variable[i]} vs {numerical_variable[j]}')
        plt.show()
```






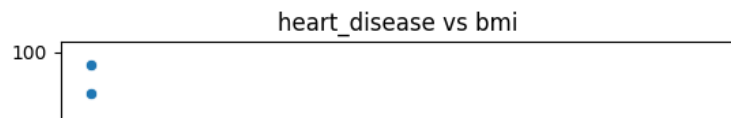
















▼ Task 7. Grouping and Aggregation:

Group data by categorical variables and calculate summary statistics within each group.

```
for Feature in categorical_variable:
    grouped = dataset.groupby(Feature)
    summary_statistics = grouped[numerical_variable].agg(['mean', 'median', 'std', 'count'])
    summary_statistics.columns = [f'{col}_{agg}' for col, agg in summary_statistics.columns]
```



```
summary_statistics.reset_index(inplace=True)
print(summary_statistics)
```

```
2      Private  36951.227009  37492.0  21257.455472  2925  45.503932
3  Self-employed  35551.288156  35315.0  20828.468932  819  60.201465
4      children  35769.432314  35106.0  21005.291797  687   6.841339

   age_median  age_std  age_count  hypertension_mean  ...  \
0      51.0    15.438879      657          0.111111  ...  
1      16.0     2.342899       22          0.000000  ...  
2      45.0    18.444200     2925          0.096068  ...  
3      63.0    16.835961      819          0.175824  ...  
4       6.0     4.533364      687          0.000000  ...  

   avg_glucose_level_std  avg_glucose_level_count  bmi_mean  bmi_median  \
0          47.697200          657  30.522063      29.40
1          28.697132           22  25.545455      23.15
2          45.886102         2925  30.304625      28.90
3          51.719490          819  30.211871      29.40
```

1 52.256627 885 30.747192 29.8




Explore differences or patterns between different groups




These conclusions provide a detailed overview of the summary statistics for each categorical variable and how they relate to various numerical fields, including age, hypertension, BMI, and stroke. They offer insights into potential patterns and trends within the dataset based on different categories.


▼ Task 8. Data Transformation:

Apply mathematical transformations (e.g., logarithmic or exponential transformations) to normalize data

```
for Feature in numerical_variable:
    dataset['log_transformed'] = np.log1p(dataset[Feature])
    plt.subplot(1, 2, 1)
    plt.hist(dataset[Feature], bins=30)
    plt.title(f'Original {Feature}')
    plt.subplot(1, 2, 2)
    plt.hist(dataset['log_transformed'], bins=30)
    plt.title('Log-Transformed Data')
    plt.show()
```







Convert categorical variables to numerical format using encoding techniques

```
from sklearn.preprocessing import LabelEncoder

label_encoder = LabelEncoder()

for Feature in categorical_variable:
    data[Feature] = label_encoder.fit_transform(dataset[Feature])

for column in dataset.select_dtypes(include='object').columns:
    label_encoder.fit(dataset[column])
    encoding_mapping = dict(zip(label_encoder.classes_, label_encoder.transform(label_encoder.classes_)))
    print(f"Encoding Mapping for {column}:")
    print(encoding_mapping)
    print()

data[categorical_variable]
```

```
Encoding Mapping for gender:
{'Female': 0, 'Male': 1, 'Other': 2}

Encoding Mapping for ever_married:
{'No': 0, 'Yes': 1}

Encoding Mapping for work_type:
{'Govt_job': 0, 'Never_worked': 1, 'Private': 2, 'Self-employed': 3, 'children': 4}

Encoding Mapping for Residence_type:
{'Rural': 0, 'Urban': 1}

Encoding Mapping for smoking_status:
{'Unknown': 0, 'formerly smoked': 1, 'never smoked': 2, 'smokes': 3}
```

	gender	ever_married	work_type	Residence_type	smoking_status
0	1	1	2	1	1
1	0	1	3	0	2
2	1	1	2	0	2
3	0	1	2	1	3

These encoding mappings are useful for machine learning and data analysis tasks as they convert categorical data into a format that can be used by algorithms that require numerical inputs. They make it easier to work with categorical data in analytical pipelines and enable the use of these variables in predictive modeling and statistical analyses.

5106	0	1	3	1	2
------	---	---	---	---	---