

PyReason

Agenda



Technical Preliminaries

The PyReason framework

Planned integration with ARL
Battlespace

Technical Preliminaries

Propositional Logic

Semantics

Implication/Rules

Fixpoint Operator

First Order Logic

Propositional Logic

Jack has school and school starts at 7 am. So Jack wakes up early.

Propositional Logic

Jack has school and school starts at 7 am. So Jack wakes up early.

A

\wedge

B

\rightarrow

C

Propositional Logic

Jack has school and school starts at 7 am. So Jack wakes up early.

A

\wedge

B

\rightarrow

C

PyReason

- **Atoms:** A, B, C, . . . (either *True* or *False*)
- **Operators:** \neg , \wedge , \vee , \rightarrow , \leftrightarrow
- **Formulas:**
 - A
 - $\neg A$
 - $A \vee B$
 - $((\neg A) \wedge B) \rightarrow C$

Semantics

- Consider a set of atoms

$$U = \{a_1, a_2, a_3\}$$

- Then we can define a world W as a subset of U

$$W = \{\}, \{a_1\}, \{a_2\}, \{a_3\}, \{a_1, a_2\}, \{a_1, a_3\}, \{a_2, a_3\}, \{a_1, a_2, a_3\}$$

Intuition: if an atom is a member of a world, it is considered true in that world otherwise it is false.

Implication / Rules

- Consider formulas: f, f', f''
- Example of a rule:

$$\frac{f' \vee f'' \rightarrow f}{\text{body} \rightarrow \text{head}(\text{atoms/negations})}$$

- Alternatively, we can write this as:

$$f \leftarrow f' \vee f''$$

- A fact is a rule with no body (i.e. body is always true)

$$f \leftarrow$$

The Fixpoint Operator (Γ):

- An application of Γ involves:
 - **Input** A set of atoms U , A world w , A set of rules R .
 - **Apply** all rules in R satisfied by w .
 - **Output** w + any atoms concluded from applied rules.

The Fixpoint Operator (Γ):

- An application of Γ involves:
 - **Input** A set of atoms U , A world w , A set of rules R .
 - **Apply** all rules in R satisfied by w .
 - **Output** w + any atoms concluded from applied rules.
- Consider,
 - $U = \{a_1, a_2, a_3\}$
 - $R = \{a_1 \leftarrow, a_2 \leftarrow a_1, a_3 \leftarrow a_2\}$
 - $w_1 = \{a_2\}$
 - $\Gamma_R(w_1) = \{a_1, a_2, a_3\}$

The Fixpoint Operator (Γ):

- Consider,
 - $U = \{a_1, a_2, a_3\}$
 - $R = \{a_1 \leftarrow, a_2 \leftarrow a_1, a_3 \leftarrow a_2\}$
 - $W_2 = \{a_3\}$
 - $\Gamma_R(W_2) = \{a_1, a_3\}$
 - $W_3 = \{\}$
 - $\Gamma_R(W_3) = \{a_1\}$

The Fixpoint Operator (Γ):

- An application of Γ involves:
 - **Input** A set of atoms U , A world w , A set of rules R .
 - **Apply** all rules in R satisfied by w .
 - **Output** w + any atoms concluded from applied rules.

- Can be written as:

$$\Gamma_R(w) = w \cup \bigcup_{r \in R} \{head(r) \text{ such that } body(r) \subseteq w\}$$

- Γ can be iteratively applied multiple times as:

$$\Gamma_R^{(i)}(w) = \Gamma_R(\Gamma_R^{(i-1)}(w))$$

The Fixpoint Operator (Γ):

- Useful for making **conclusions**, as well as, **explanations** behind them:

- $U = \{a_1, a_2, a_3\}$
- $R = \{a_1 \leftarrow, a_2 \leftarrow a_1, a_3 \leftarrow a_2\}$

- $W_3 = \{\}$
- $\Gamma_R^{(1)}(W_3) = \{a_1\}$
- $\Gamma_R^{(2)}(W_3) = \{a_1, a_2\}$
- $\Gamma_R^{(3)}(W_3) = \{a_1, a_2, a_3\}$

$a_1 \leftarrow$

$a_2 \leftarrow a_1$

$a_3 \leftarrow a_2$

First Order Logic/Predicate Calculus

Predicates are a way to specify atomic propositions.

Consider,

“friend” is a predicate

v_1, v_2 are two variables

$\text{friend}(v_1, v_2)$

First Order Logic/Predicate Calculus

Predicates are a way to specify atomic propositions.

Consider,

“friend” is a predicate

v_1 , v_2 are two variables

jack, phil are two people (constants)

friend(v_1 , v_2)



friend(jack, phil)

First Order Logic/Predicate Calculus

Predicates are a way to specify atomic propositions.

Consider,

“friend” is a predicate

v_1, v_2 are two variables

jack, phil are two people (constants)

friend(v_1, v_2)



friend(jack, phil)



Jack and Phil are friends

First Order Logic/Predicate Calculus

Predicates are a way to specify atomic propositions.

Consider,

“friend” is a predicate

v_1 , v_2 are two variables

jack, phil are two people (constants)

friend(v_1 , v_2)



friend(jack, phil)



Jack and Phil are friends

First Order Logic/Predicate Calculus

Non-ground atoms are the key item that differentiates Predicate Calculus from Propositional Calculus.

Predicate + Variable symbol(s) =
(Non-ground) atomic proposition

Predicate + Constant(s) =
(Ground) atomic proposition

friend(v_1 , v_2)




Grounding

friend(jack, phil)

Predicate Calculus in Knowledge Graphs

- Unary predicates can model attributes of nodes.
e.g. $\text{student}(c_1)$
- Binary predicates can model relationships between nodes (attributes of edges).
e.g. $\text{friend}(c_1, c_2)$



The PyReason framework

Lattice structure and annotations

Support for Temporal Reasoning

Notion of Interpretation

Rules

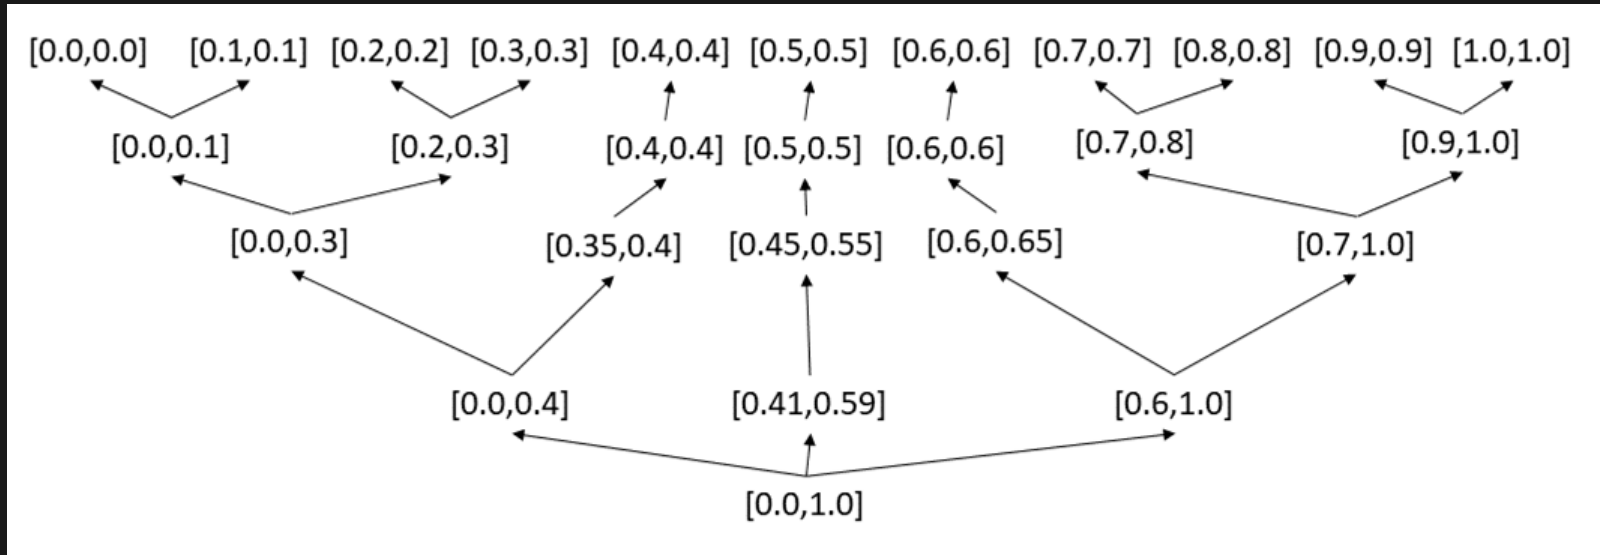
Type-checking and Consistency checking

Generalized Annotated Logic

PyReason performs reasoning about first-order and propositional logic statements,

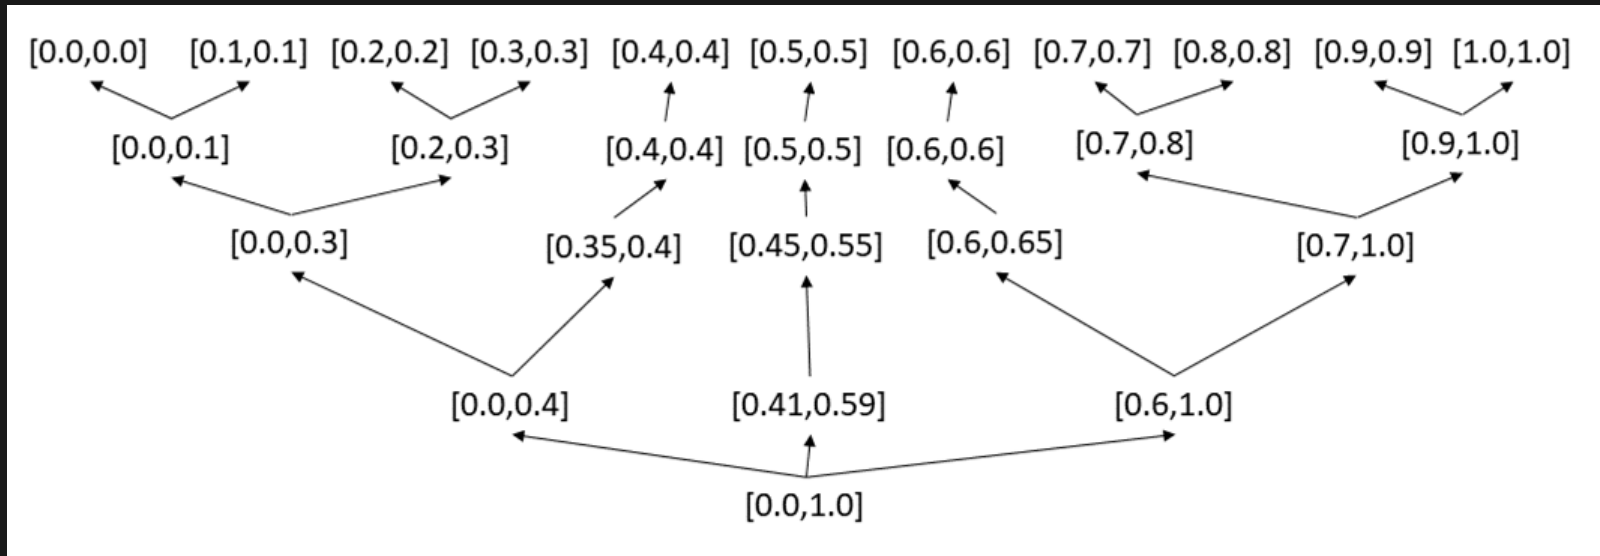
Generalized Annotated Logic

PyReason performs reasoning about first-order and propositional logic statements, that can be annotated with elements of a lattice structure.



Generalized Annotated Logic

PyReason performs reasoning about first-order and propositional logic statements, that can be annotated with elements of a lattice structure.



Interpretations map atomic propositions to elements of this lattice structure

Design Feature: Uncertainty

Allowing Interpretation (I) to map atoms to bounds allows us to model uncertainty effectively.

For e.g.

When we have no information about `friend(jack,phil)`,

$$I(\text{friend}(\text{jack}, \text{phil})) = [0, 1]$$

While still supporting the propositional cases:

$$I(\text{friend}(\text{jack}, \text{phil})) = [1, 1]$$

friends

$$I(\text{friend}(\text{jack}, \text{phil})) = [0, 0]$$

not friends

Design Feature: Temporal Reasoning

We additionally allow Interpretation (I) to map time (alongside atoms) to bounds and hence we can perform reasoning over time.

Continuing with the same example, we can have,

$$I(\text{friend}(\text{jack}, \text{phil}), \text{jan}) = [0, 0]$$

not friends in January

$$I(\text{friend}(\text{jack}, \text{phil}), \text{feb}) = [0, 1]$$

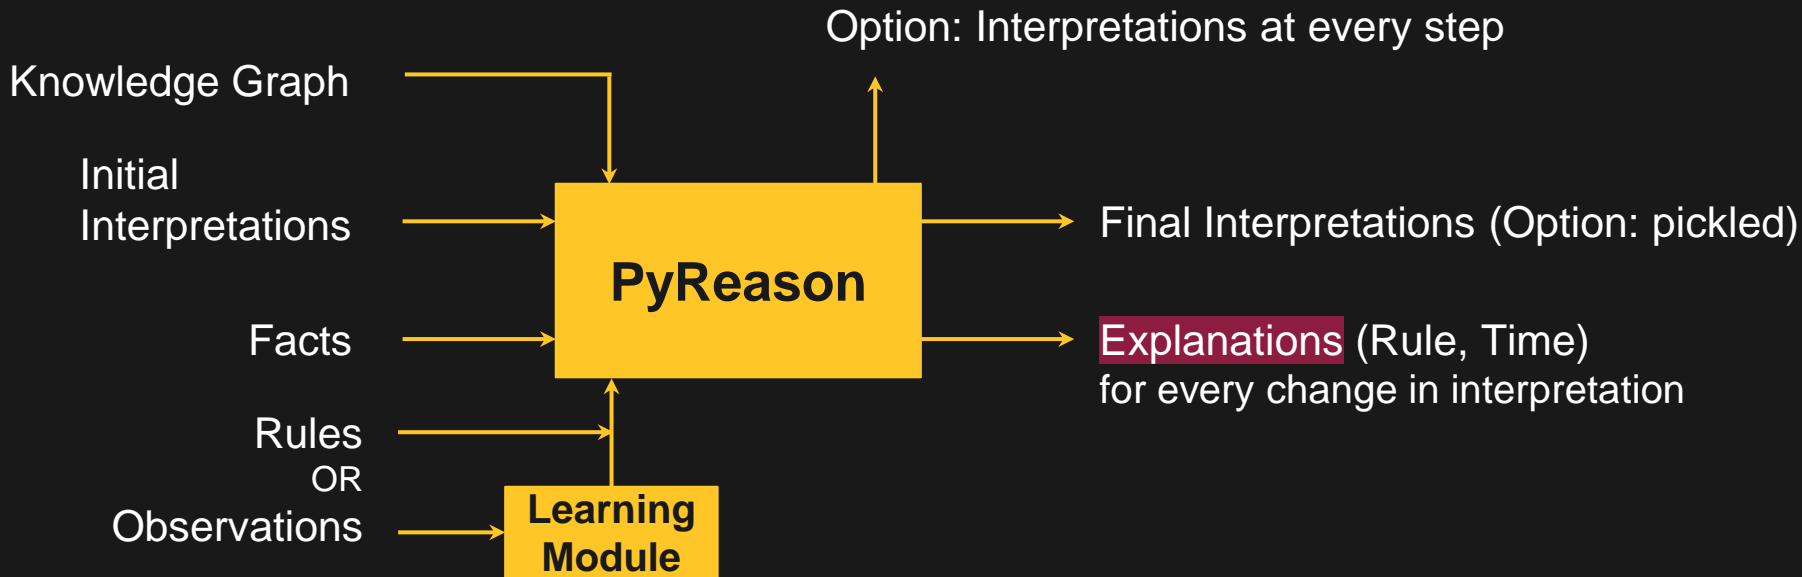
no info about February

open-world assumption

$$I(\text{friend}(\text{jack}, \text{phil}), \text{mar}) = [1, 1]$$

friends in March

PyReason Input/Outputs



Design Feature: Explainability

By implementing the fixpoint operator directly (as opposed to a black box heuristic) the software enables full explainability of the result.

- We can recover a trace of every rule applied and its effect.
- We can uncover causal relationships between atomic propositions.
- We can detect logical flaws and inconsistencies.

Logical Rules Reasoning within a node

Universally quantified non-ground rule

$$\boxed{\forall X:} \text{pred}(X): \boxed{f(x_1, \dots, x_n)} \leftarrow_{\Delta t} \bigwedge_{\text{pred}_i \in \text{UnaSet}} \text{pred}_i(X): x_i$$

Universal quantifier
(*design feature*)

Annotation is a function over the elements in lattice
e.g. *Max, Min, Avg,*

Fuzzy t-norms and conorms
(*design feature*)

Logical Rules Reasoning across an edge

Universally quantified non-ground rule

$$\forall X: pred(X): f(x_1, \dots, x_n) \leftarrow_{\Delta t} \boxed{\exists_k} X': \boxed{rel(X, X'): [1,1]} \wedge \bigwedge_{pred_q \in BinSet} pred_q(X, X'): x_q$$
$$\wedge \bigwedge_{pred_r \in Unaset} pred_r(X): x_r \wedge \bigwedge_{pred_s \in Unaset'} pred_s(X'): x_s$$

Existential quantifier
(*design feature*)

'rel' is a reserved word.

Marked portion denotes that -
an edge exists between X and X'

Logical Rules Examples

- $\text{promoted}(X): [1,1] \leftarrow_{1\text{year}} \text{gpa}(X): [0.3, 1]$
Student X will get promoted at the end of the year if their overall gpa is in the top 70% of the class.
- $\text{promoted}(X): [1,1] \leftarrow_{1\text{year}} \text{gpa}(X): [0.3, 1] \wedge \wedge_Y \text{takes}(X, Y): [1,1] \wedge \wedge_Y \text{passed}(X, Y): [1,1]$
adds an additional condition that to get promoted, X must pass all of the courses they take.
- $\text{gpa}(X): [\text{avg}(x_s)] \leftarrow_{1\text{year}} \exists_2 \text{takes}(X, Y): [1,1] \wedge \text{score}(X, Y): x_s$
shows a way to compute gpa using two classes taken by X.

Design Feature: Type-checking

student(eve) ✓

student(cal) ✓

subject(math) ✓

friend(eve, cal) ✓

takes(eve, math) ✓

student(math) ✗


subject(eve) ✗

subject(cal) ✗

friend(math, cal) ✗


takes(eve, cal) ✗

takes(math, eve) ✗




Design Feature: Type-checking

- Avoids silly errors like: *“Math is driving a car”*.
- Provides drastic reduction to complexity induced by the grounding problem, by increasing graph sparsity, reducing storage and computations.
- Significantly improves utility in a variety of application domains.



Design Feature: Type-checking

- Optional feature. Can be turned on/off.
- Specified at the time of building the graph: if we have prior knowledge about constraints over predicate-atom pairs.



Design Feature: Support for literals

A literal is any ground atom or a negation of a ground atom.

Traditional logic frameworks only support atoms or its negation, not both.

In PyReason -

- Atoms and their negations can be **simultaneously implemented**.
- We define both as separate ground atoms, and,
- We define a **consistency constraint** that prevents an atom and its negation to co-exist.

Design Feature: Consistency for pairs

Defining literals:

$I(\text{at_home}(x),t) = [1,1]$ and $I(\text{not_home}(x),t) = [1,1]$ X

Modelling relationships between pairs which might become inconsistent:

$I(\text{bachelor}(x),t) = [1,1]$ and $I(\text{married}(x),t) = [1,1]$ X

$I(\text{fit}(x),t) = [0.6,1]$ and $I(\text{injured}(x),t) = [0,0.8]$ X

Design Feature: Consistency for pairs

Checking:

$$I_1 = [L_1, U_1]$$

$$I_2 = [L_2, U_2]$$

Then they are consistent,

$$\text{Iff } L_1 \leq 1 - L_2 \quad \text{and} \quad U_1 \geq 1 - U_2$$

Resolution:

$$I_1 = I_2 = [0, 1]$$

Design Feature: Consistency during execution

Checking:

$$I_{\text{current}} = [L_1, U_1]$$

$$I_{\text{new}} = [L_2, U_2]$$

Then they are consistent,

$$\text{Iff } L_1 \leq U_2 \quad \text{and} \quad U_1 \geq L_2 \quad \text{and} \quad L_2 \leq U_2$$

from outcome of a rule

i.e. from same lower lattice

Resolution:

$$I_{\text{updated}} = [0, 1]$$

Integration within ARL Battlespace

Modelling the Game World

Interfacing with PyReason

Modeling the Game World

- The game board is modeled as a graph:
 - Each square is a node.
 - Edges between neighbouring squares, air and land.

Modeling the Game World

- The game board is modeled as a graph:
 - Each **square** is a node.
 - Edges between **neighbouring squares**, air and land.
- Other components can be modeled as -
 - **Units** (Soldier, Tank, Truck, Flag, Airplane, Missiles) are nodes, with attribute '**type**' = {air, ground, immovable}.
 - **Players** (A,B,C,D) are nodes.

Modeling the Game World

- The game board is modeled as a graph:
 - Each **square** is a node.
 - Edges between **neighbouring squares**, air and land.
- Other components can be modeled as -
 - **Units** (Soldier, Tank, Truck, Flag, Airplane, Missiles) are nodes, with attribute '**type**' = {air, ground, immovable}.
 - **Players** (A,B,C,D) are nodes.
 - Edges between Units and Squares have **at(U,S)**.
 - Edges between Players and Units have **of(P,U)**.

...

Modeling the Game World

- Rules:

- All possible actions

e.g. Advancement changes location of an unit

e.g. Rotate changes orientation.

Modeling the Game World

- Rules:

- All possible **actions**

e.g. Advancement changes location of an unit

- Movement of missiles are **facts**

Modeling the Game World

- Rules:

- All possible **actions**
- Movement of missiles are **facts**
- Causal effects of actions

e.g. Advancement changes location of an unit

e.g. Overlap leads to mutual destruction

Advance U1, S1

Advance U2, S1

Trigger: Destroy U1, U2

Modeling the Game World

- Rules:
 - All possible **actions** *e.g. Advancement changes location of an unit*
 - Movement of missiles are **facts**
 - Causal effects of actions *e.g. Overlap leads to mutual destruction*
 - . . .
- **Termination** conditions (flag capture / annihilation) is captured in the body of a rule, which when fired ends the game.

Modeling the Game World

- Rules:
 - All possible **actions** *e.g. Advancement changes location of an unit*
 - Movement of missiles are **facts**
 - Causal effects of actions *e.g. Overlap leads to mutual destruction*
 - . . .
- **Termination** conditions (flag capture / annihilation) is captured in the body of a rule, which when fired ends the game.
- At initialization, type-checking ensures attributes are matched to appropriate nodes and edges in the graph.

Interfacing with PyReason

Input: Current State

List of interpretations with bounds in .yaml format

Course of Action

List of tuples of the form (action, player, unit, time)

Output: Next States

List of interpretations with bounds in .pkl format



Thank You

V2
LAB