

PRIMA: Planner- Reasoner Inside a Multi- task Reasoning Agent

Daoming Lyu Auburn University, Auburn,

Bo Liu Auburn University, Auburn,

13 Feb 2022

Submitted to

Sir Muhammad Ata Ur
Rehman

Group Members

- Munim Siddiqui
(20K-1730)
- Murtaza Salman
(20K-047)
- Ali Qureshi
(20K-0149)

Contents



Introduction & Background

Problem Formulation & Limitations

PRIMA: Planner-Reasoner for Multi-task Reasoning

An MDP formulation of the learning-to-reason problem

Experimental Results and Analysis

Conclusion

Introduction to Multi-task Learning (MTL)

- **MTL** is an approach that improves sample complexity by learning to solve multiple tasks simultaneously, as opposed to the traditional "one model per task" strategy.
- Enhances AI's efficiency and predictive accuracy by leveraging the commonalities across tasks
- A significant challenge of MTL is the management of its model size that is larger than the single-task models.
- **Machine Task Learning** struggles with large models due to the conflict between generalization and task-specific efficiency.
 - **Generalization:** refers to a model's ability to perform well on unseen data i.e data not used during training.
 - **Task-Specific Efficiency:** Concerns how well a model can perform a specific task
- **PRIMA architecture**, trained using deep reinforcement learning, achieves a balance between capability and efficiency in MTL, yielding state-of-the-art performance.

Background

- **MTL Evolution:** Demonstrates improved generalization by leveraging shared knowledge across tasks.
- **Deep Learning Integration:** Incorporates deep neural networks to enhance MTL's success and applicability.
- **Inherent Challenges:** Balances model size and task-specific efficiency while preserving generalization.
- **Logic Reasoning:** Employs First-Order Logic for deducing new information from known facts in AI.
- **PRIMA Framework:** Introduces a approach to efficiently balance reasoning capabilities in MTL.




Figure 1: An example from the AdjacentToRed task and its formulation as a logical reasoning problem.

Problem Formulation

- **Simplification of First-Order Logic (FOL):** Focuses on variables, constants, and predicates to define relationships as True or False.
- **Construction of Logical Formulas:** Describes the construction of logical expressions using atoms, connectives, and quantifiers.
- **Definition of Multi-task Reasoning (MTR):** Aims to predict outputs across tasks, emphasizing joint learning of multiple reasoning tasks.
- **Objective of MTR:** Seeks to enhance task performance by applying reasoning skills across various tasks in a unified model.

Limitations of Current Multi-task Learning Models

- **Generalization vs. Efficiency Conflict:** Current MTL models struggle to balance broad generalization across tasks with the need for efficient, task-specific execution.
- **Model Size Dilemma:** The large size of conventional MTL models poses challenges which can slow them down and make them less effective at specific tasks..
- **Capability Overreach:** By trying to do too much, MTL models might not do any one task particularly well.
- **Not flexible enough:** These models can't easily adjust to the specific needs of each task, which can make them less efficient overall.

3. Model Architecture of PRIMA

- PRIMA helps us reason about multiple things at once.
- We use a strategy called forward-chaining to figure out answers step by step.
- But it's too slow to consider every possibility, so we need to be smarter about how we do it.
- PRIMA has two main parts: the Reasoner and the Planner.
- The Reasoner figures out logical rules using neural networks.
- The Planner decides which rules to use and how to put them together.
- We train PRIMA to be good at reasoning efficiently using deep reinforcement learning.

3.1 Reasoner: transforming logic rules into neural operators

- The Reasoner is like a smart assistant that understands logic and helps us find answers.
- It uses neural networks to translate logical rules into something it can understand.
- It focuses on three main rules: Boolean Logic, Expansion, and Reduction, which help it make sense of logical statements.

3.2 Planner: activating and forward-chaining learnable Horn clauses

- The Planner is like a strategist that plans how to solve problems efficiently.
- It decides which logical rules to use and in what order, sort of like picking the best path through a maze.
- Instead of trying every possible combination, it's smart and picks the most promising routes.
- This way, it finds solutions faster without getting stuck in unnecessary details.

4. Learning-to-Reason via Reinforcement Learning

- We're using reinforcement learning (RL) to teach a computer to reason logically.
- The Planner makes decisions about which neural operators to use based on binary indicators.
- This process is similar to making sequential decisions in Markov decision processes (MDPs)

4.1 An MDP formulation of the learning-to-reason problem

- **MDP Definition:** Think of an MDP as a framework with states, actions, transition probabilities, rewards, and a discount factor.
- **In Our Context:** The states represent sets of predicates, actions are binary indicators for activating neural operators, and rewards are based on reasoning accuracy.
- **Objective:** We want to find the best policy (sequence of actions) that maximizes the total expected reward over a series of reasoning steps.

4.2 Policy network: modeling of planner

- **Planner Module:** It's like the brain deciding what to do. It's represented by a policy network.
- **Base Planner:** It's a simple version of the Planner, deciding which operators to activate based on input predicates.
- **Enhanced Planner:** This version uses Monte Carlo Tree Search (MCTS) to improve decision-making by exploring possible reasoning paths.

4.3 Overall learning framework

- We're using RL methods to teach the system to reason. We'll try REINFORCE and MuZero.
- **REINFORCE:** It's a model-free RL method, meaning it learns directly from experiences without a model of the environment.
- **MuZero:** It's a model-based RL method that combines planning and learning, and it's good for complex tasks.
- We use MuZero because it handles partial structural knowledge well, reducing testing complexity.
- **Training Process:** We load the model weights, execute reasoning path rollouts according to the policy network, and then train the Planner-Reasoner using the data collected.

5 Experimental Results and Analysis

- Evaluate the performance of different variants of PRIMA on eight tasks from the family tree and graph benchmarks.
- Tasks are widely used benchmarks for inductive logic programming.
- Evaluate their testing accuracy and reasoning cost

Outdegree

Adjacent To Red

HasFather

HasSister

Connectivity

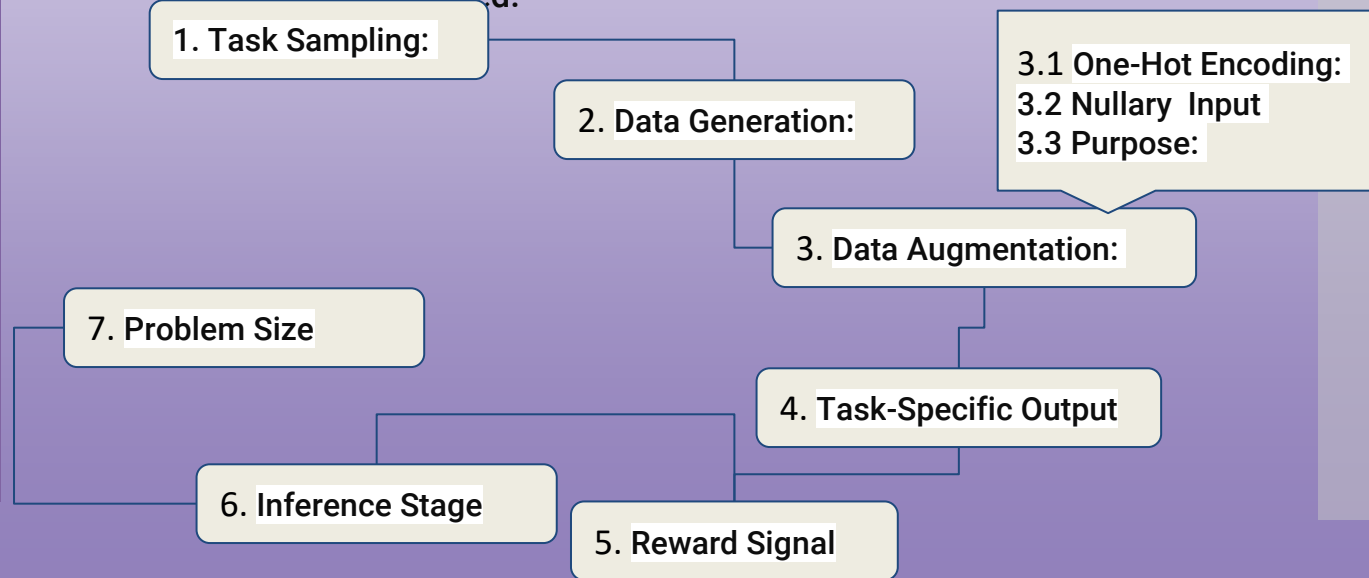
Is Grandparent

IsUncle

IsMGUncle

5.1 Experimental Setups

- Methodology used for training and testing in the multi-task setting, particularly focusing on how tasks are sampled, data is generated, and the model is trained and evaluated.



5.2 Overall performance


- **Single-task reasoning capability**
- Before we proceed to evaluate the MTR capabilities of PRIMA, We adapt to the single-task setting and train a separate model for each task. We name this

testing accuracy			Family Tree	HasFather	HasSister	IsGrandparent	IsUncle	IsMGUncle	Graph	AdjacentToRed	4-Connectivity	1-OutDegree
	Single Task	PRISA-REINFORCE	m=20	62.6	50.7	96.5	97.3	99.8	m=10	47.7	33.5	48.7
			m=100	87.8	69.8	2.3	97.7	98.4	m=50	71.6	92.8	97.4
			PSS	0	0	0	0	0	PSS	0	0	0
		PRISA-PPO	m=20	71.5	64.3	97.5	98.1	99.6	m=10	62.3	57.8	61.6
			m=100	93.2	78.7	98.2	97.3	99.1	m=50	85.5	95.2	96.3
			PSS	0	0	0	0	0	PSS	0	0	0
		PRISA-MuZero	m=20	100	100	100	100	100	m=10	100	100	100
			m=100	100	100	100	100	100	m=50	100	100	100
			PSS	100	100	100	100	100	PSS	90	100	100

5.2 Overall performance

- Multi-task reasoning capability

		Family Tree	HasFather	HasSister	IsGrandparent	IsUncle	IsMGUncle	Graph	AdjacentToRed	4-Connectivity	1-OutDegree	
testing accuracy	Single Task	MemNN	m=20	99.9	86.3	96.5	96.3	99.7	m=10	95.2	92.3	99.8
			m=100	59.8	59.8	97.7	96	98.4	m=50	93.1	81.3	78.6
			PSS	0	0	0	0	0	PSS	0	0	0
	Multi-Task	∂ ILP/ NLM/ DLM	m=20	100	100	100	100	100	m=10	100	100	100
			m=100	100	100	100	100	100	m=50	100	100	100
			PSS	100	100	100	100	100/ 90/ 100	100/ 20/ 70	PSS	100/ 90/ 90	100
		DLM-MTR	m=20	100	100	100	100	100	m=10	96.7	100	100
			m=100	100	100	100	100	100	m=50	97.2	100	100
			PSS	100	100	100	100	100	PSS	0	100	100
		PRIMA	m=20	100	100	100	100	100	m=10	100	100	100
			m=100	100	100	100	100	100	m=50	100	100	100
			PSS	100	100	100	100	100	90	PSS	90	100



6. Conclusion

- long-standing challenge in multi-task reasoning (MTR) is the intrinsic conflict between capability and efficiency
- development of a novel neural-logic architecture termed PRIMA
- PRIMA improves inference efficiency and achieves a state-of-the-art balance between MTR capability and inference efficiency.
- The model's training follows a complete data-driven end-to-end approach via deep reinforcement learning, and the performance is validated across a variety of benchmarks.

ANY
QUESTIONS?

