The background features abstract, overlapping geometric shapes in various shades of blue, ranging from light sky blue to deep navy blue. These shapes are primarily located on the left and right sides of the slide, framing the central text.

# Chapter 21

## Database Recovery Techniques

# Recovery Concepts

- ▶ Recovery Outline and Categorization of Recovery Algorithms
- ▶ Recovery from transaction failures:
  - ▶ the database is **restored** to the most recent consistent state before the time of failure.
  - ▶ The system keeps information about the changes that were applied to data items by the various transactions in the **system log**.
  - ▶ In case of failure, Just recover the committed transactions as uncommitted transactions doesn't modify database. So, they will be rolled back

# Recovery Concepts

- ▶ Recovery Outline and Categorization of Recovery Algorithms
- ▶ **CATASTROPHIC FAILURE:**
- ▶ Extensive damage to database because of disk crash,
- ▶ **Solution:**
  - ▶ **restores a past copy** of the db. from some archival storage
  - ▶ and **redo the operations of committed transactions** from the log file up to failure point to ensure Database consistency.

# Recovery Concepts

## ▶ Recovery Outline and Categorization of Recovery Algorithms

### ▶ NON-CATASTROPHIC FAILURE:

- ▶ When the database on disk is not physically damaged,
  - ▶ the recovery strategy is **to identify any changes** that may cause an inconsistency in the database.
- ▶ For example,
  - ▶ A transaction has not yet committed and is failed,
  - ▶ Then it is possible that some of its write operations are done and some are left.
  - ▶ So all the operations needs to be undone
  - ▶ And redo for the whole transaction is required to ensure atomicity & data consistency

# Recovery Concepts

- ▶ Recovery Outline and Categorization of Recovery Algorithms
- ▶ **Recovery from non-Catastrophic failure:**
- ▶ Conceptually, we can define 2 policies for recovery from non-catastrophic failures:
  - ▶ Deferred update
  - ▶ Immediate update

# Recovery Concepts

- ▶ Recovery Outline and Categorization of Recovery Algorithms
- ▶ Deferred update techniques:
- ▶ Do not update the database on disk before COMMIT operation
- ▶ Before COMMIT,
  - ▶ all changes done by the transactions are saved in local workspace or the Main memory buffers.
  - ▶ all the updates operations are recorded in log file, once COMMIT is done, all the updates are written to database

# Recovery Concepts

- ▶ Recovery Outline and Categorization of Recovery Algorithms
- ▶ Deferred update techniques:
- ▶ In case of transaction failure:
  - ▶ No need to UNDO the changes done by the transaction as the changes were done locally
  - ▶ Just REDO of the operations is required as the changes were not written to database yet.
  - ▶ That's why Deferred update is also known as the **NO-UNDO/REDO** algorithm.

# Recovery Concepts

- ▶ Recovery Outline and Categorization of Recovery Algorithms
- ▶ Deferred update techniques:
- ▶ Example: (Log For Deferred Update)
  - ▶ <start\_Transaction,T1>
  - ▶ <read\_item,T1,A>
  - ▶ <write\_item,T1,A,10>
  - ▶ <commit, T1>
- ▶ Note: here, we don't need to store outdated value of data item because in this case we never undo any operation



# Recovery Concepts

- ▶ Recovery Outline and Categorization of Recovery Algorithms
- ▶ Immediate update technique:
- ▶ The database is updated by some operations of a transaction before reaching the COMMIT point
  - ▶ i.e., Writing in main memory and disk side by side without waiting for COMMIT operation
  - ▶ All operations are forcefully written in log file before COMMIT to make recovery possible.
- ▶ If a transaction fails after some operations but before COMMIT
  - ▶ the effect of its operations on the database must be undone.
- ▶ This technique, known as the **UNDO/REDO** algorithm, requires both operations during recovery

# Recovery Concepts

- ▶ Recovery Outline and Categorization of Recovery Algorithms:
- ▶ Immediate update technique
- ▶ Example: (Log For ImmediateUpdate)
  - ▶ <start\_Transaction,T1>
  - ▶ <read\_item,T1,A>
  - ▶ <write\_item,T1,A,10,11>
  - ▶ <commit, T1>
  - ▶ Note: here, we store old and new values of data item

# Recovery Concepts

- ▶ **Caching (Buffering) of Disk Blocks**
- ▶ For recovery purposes, buffering of disk blocks is done in DBMS cache
- ▶ DBMS cache:
  - ▶ Collection of in-memory buffers.
- ▶ A **directory** for the cache is maintained to track which **database items are in the buffers**.
- ▶ Directory is a table that saves entries as:  
**<Disk\_page\_address, Buffer\_location, ... >**

# Recovery Concepts

- ▶ **Caching (Buffering) of Disk Blocks**
- ▶ Whenever, DBMS requests some item:
  - ▶ cache directory is checked whether the requested item is in DBMS cache or not.
  - ▶ If it is not, the item must be located on disk, So the appropriate disk pages are copied into the cache.
- ▶ **If DBMS cache is already filled?**
  - ▶ cache buffers needed to be replaced to make space available for the new item

# Recovery Concepts

- ▶ Caching (Buffering) of Disk Blocks
- ▶ **The entries in DBMS cache Dictionary:**
- ▶ Holds additional info. For buffer management
  - ▶ Dirty bit
  - ▶ pin-unpin bit

# Recovery Concepts

- ▶ **Caching (Buffering) of Disk Blocks**
- ▶ **The entries in DBMS cache Dictionary:**
- ▶ **Dirty bit:**
  - ▶ Associated with each buffer in the cache
  - ▶ included in the directory entry
  - ▶ to indicate **whether or not the buffer has been modified.**
- ▶ **Dirty bit = 0 :**
  - ▶ page is first read from the database disk into a cache buffer,
  - ▶ a new entry is inserted in the cache directory with the new disk page address.
- ▶ **Dirty bit = 1 :**
  - ▶ As soon as the buffer is modified.
  - ▶ Data needs to be written to disk when the buffer contents are replaced from the cache

# Recovery Concepts

- ▶ **Caching (Buffering) of Disk Blocks**

- ▶ **Pin-unpin bit:**

- ▶ page in the cache is pinned (bit value 1) if it cannot be written back to disk as yet.
- ▶ Why?
  - ▶ Because the recovery protocol may restrict certain buffer pages from being written back to the disk until the transactions that changed this buffer have committed.

# Recovery Concepts

- ▶ Caching (Buffering) of Disk Blocks
- ▶ Strategies for replacing cache buffers
  - ▶ **In-place updating**
    - ▶ Writes the buffer to the same original disk location
    - ▶ Overwrites old values of any changed data items
  - ▶ **Shadowing**
    - ▶ Writes an updated buffer at a different disk location, to maintain multiple versions of data items
    - ▶ Not typically used in practice



# Recovery Concepts

- ▶ Caching (Buffering) of Disk Blocks
- ▶ **Before-image(BFIM):**
  - ▶ old value of data item before updating
- ▶ **After-image(AFIM):**
  - ▶ new value of data item after updating

# Recovery Concepts

- ▶ Write-Ahead Logging, Steal/No-Steal, and Force/No-Force
- ▶ **REDO-type log entry includes the new value** (AFIM) of the item written by the operation since this is needed to redo the effect of the operation from the log.
- ▶ **UNDO-type log entries include the old value** (BFIM) of the item since this is needed to undo the effect of the operation from the log.
- ▶ In an **UNDO/REDO algorithm**, both **BFIM** and **AFIM** are recorded into a single log entry.

# Recovery Concepts

- ▶ **Write-Ahead Logging, Steal/No-Steal, and Force/No-Force**
- ▶ DBMS cache holds the cached database disk blocks in main memory buffers.
- ▶ When an update to a data block—stored in the DBMS cache—is made, an associated log record is written to the log buffer in the DBMS cache.
- ▶ **Write-ahead logging approach:** the log buffers (blocks) that contain the associated log records for a particular data block update must first be written to disk before the data block itself can be written back to disk from its main memory buffer.

# Recovery Concepts

- ▶ **Write-Ahead Logging, Steal/No-Steal, and Force/No-Force**
- ▶ **No-steal approach**: a cache buffer page updated by a transaction cannot be written to disk before the transaction commits.
  - ▶ UNDO will never be needed during recovery, since a committed transaction will not have any of its updates on disk before it commits.
- ▶ **pin-unpin bit will be set to 1 (pin)** to indicate that a cache buffer cannot be written back to disk.
- ▶ **Steal**: the recovery protocol allows writing an updated buffer before the transaction commits.

# Recovery Concepts

- ▶ **Write-Ahead Logging, Steal/No-Steal, and Force/No-Force**
- ▶ **Force approach:** If all pages updated by a transaction are immediately written to disk before the transaction commits.
  - ▶ REDO will never be needed during recovery, since any committed transaction will have all its updates on disk before it is committed.
- ▶ Otherwise, it is called no-force.
- ▶ For frequently changed objects, a no-force policy reduces the number of write operations to the on-disk database object.

# Recovery Concepts

- ▶ **Write-Ahead Logging, Steal/No-Steal, and Force/No-Force**
- ▶ The deferred update (NO-UNDO) recovery scheme follows a **no-steal approach**.
- ▶ Typical database systems employ a **steal/no-force (UNDO/REDO)** strategy.
- ▶ The advantage of steal is that it **avoids the need for a very large buffer space to store all updated pages in memory**.
- ▶ The advantage of no-force is that an updated page of a committed transaction may still be in the buffer when another transaction needs to update it.

# Recovery Concepts

- ▶ **Write-Ahead Logging, Steal/No-Steal, and Force/No-Force**
- ▶ Write-ahead logging (WAL) protocol for a recovery algorithm that requires both UNDO and REDO:
- ▶ The before image of an item cannot be overwritten by its after image in the database on disk until all UNDO-type log entries for the updating transaction have been force-written to disk.
- ▶ The commit operation of a transaction cannot be completed until all the REDO-type and UNDO-type log records for that transaction have been force written to disk.

# Recovery Concepts

- ▶ Checkpoints in the System Log and Fuzzy Checkpointing
- ▶ Another type of entry in the system log is called a checkpoint.
- ▶ A checkpoint is written in log file as :
  - ▶ **[checkpoint, list of active transactions]**
- ▶ When system writes the modified DBMS buffers on disk, then this checkpoint entry is inserted into system log
- ▶ All transactions that have their [commit, T] entries in the log before a [checkpoint] entry
  - ▶ do not need to have their WRITE operations redone in case of a system crash,
  - ▶ since all their updates will be recorded in the database on disk during checkpointing.
- ▶ List of active transactions from checkpoint entry will help in identifying transactions for recovery.



# Recovery Concepts

- ▶ Checkpoints in the System Log and Fuzzy Checkpointing
- ▶ The DB recovery manager will decide **at what time the checkpoint is taken.**
  - ▶ The interval for checkpoints can be decided in terms of :
    - ▶ time (i.e., after every 4 minutes) or
    - ▶ after the number of committed transactions since the last checkpoint (i.e., after every 4 transactions' commit)

# Recovery Concepts

- ▶ Checkpoints in the System Log and Fuzzy Checkpointing
- ▶ Taking a check point involves following actions:
  - ▶ Suspend execution of transactions temporarily.
  - ▶ Force-write all main memory buffers that have been modified to disk.
  - ▶ Write a [checkpoint] record to the log, and force-write the log to disk.
  - ▶ Resume executing transactions.

# Recovery Concepts

- ▶ **Checkpoints in the System Log and Fuzzy Checkpointing**
- ▶ The cost of writing modified buffers back to disk (i.e., step 2) will induce delay in currently executing transactions as they are at halt state.
- ▶ That's why fuzzy checkpointing is used.
- ▶ **Fuzzy checkpointing.**
- ▶ The system can resume transaction processing after a [begin\_checkpoint] record is written to the log without having to wait for step 2 to finish.
- ▶ When step 2 is completed, an [end\_checkpoint, ... ] record is written in the log with the relevant information collected during checkpointing

# Recovery Concepts

- ▶ Transaction Rollback and Cascading Rollback
- ▶ If a transaction fails for some reason after updating the database, but before the transaction commits,
  - ▶ it may be necessary to roll back the transaction.
- ▶ If any data item values have been changed by the transaction and written to the database on disk,
  - ▶ they must be restored to their previous values (BFIMs).
- ▶ The **undo-type log entries** are used to restore the old values of data items that must be rolled back.
  - ▶ Because they contains both old and new values

# Recovery Concepts

- ▶ **Transaction Rollback and Cascading Rollback**
- ▶ **cascading rollback:**
- ▶ If a transaction  $T$  is rolled back, any transaction  $S$  that has read the value of some data item  $X$  written by  $T$ , must also be rolled back.
- ▶ Similarly, once  $S$  is rolled back, any transaction  $R$  that has read the value of some data item  $Y$  written by  $S$  must also be rolled back; and so on.
- ▶ it can occur when the recovery protocol ensures recoverable schedules but does not ensure strict or cascadeless schedules.

# Recovery Concepts

## ► Transaction Rollback & Cascading Rollback

$T_1$
read_item( $A$ )
read_item( $D$ )
write_item( $D$ )

$T_2$
read_item( $B$ )
write_item( $B$ )
read_item( $D$ )
write_item( $D$ )

$T_3$
read_item( $C$ )
write_item( $B$ )
read_item( $A$ )
write_item( $A$ )

# Recovery Concepts

		<i>A</i>	<i>B</i>	<i>C</i>	<i>D</i>
		30	15	40	20
	[start_transaction, $T_3$ ]				
	[read_item, $T_3$ , <i>C</i> ]				
*	[write_item, $T_3$ , <i>B</i> , 15, 12]		12		
	[start_transaction, $T_2$ ]				
	[read_item, $T_2$ , <i>B</i> ]				
**	[write_item, $T_2$ , <i>B</i> , 12, 18]		18		
	[start_transaction, $T_1$ ]				
	[read_item, $T_1$ , <i>A</i> ]				
	[read_item, $T_1$ , <i>D</i> ]				
	[write_item, $T_1$ , <i>D</i> , 20, 25]				25
	[read_item, $T_2$ , <i>D</i> ]				
**	[write_item, $T_2$ , <i>D</i> , 25, 26]				26
	[read_item, $T_3$ , <i>A</i> ]				

← System crash

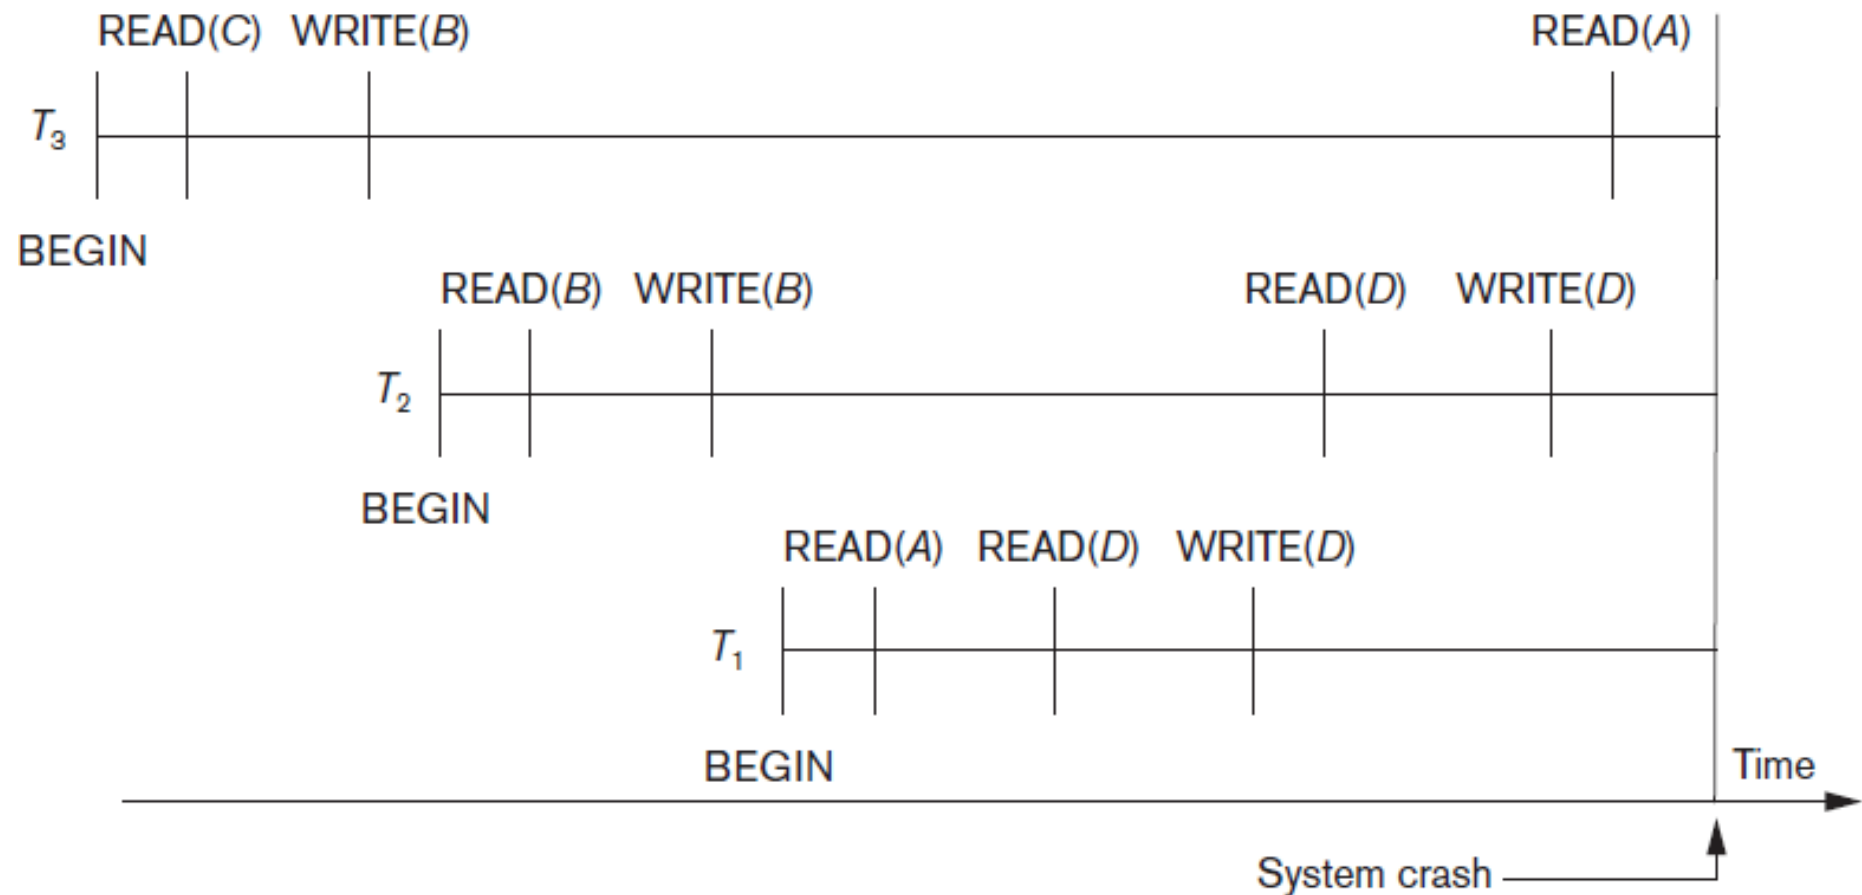
$T_1$	$T_2$	$T_3$
read_item( <i>A</i> )	read_item( <i>B</i> )	read_item( <i>C</i> )
read_item( <i>D</i> )	write_item( <i>B</i> )	write_item( <i>B</i> )
write_item( <i>D</i> )	read_item( <i>D</i> )	read_item( <i>A</i> )
	write_item( <i>D</i> )	write_item( <i>A</i> )

\*  $T_3$  is rolled back because it did not reach its commit point.

\*\*  $T_2$  is rolled back because it reads the value of item *B* written by  $T_3$ .

# Recovery Concepts

## ► Transaction Rollback and Cascading Rollback





# Recovery Concepts

## ► Transaction Rollback and Cascading Rollback

We must now check for cascading rollback. T2 reads the value of item B that was written by transaction T3. Because T3 is rolled back, T2 must now be rolled back, too. The WRITE operations of T2, marked by \*\* in the log, are the one that are undone.

Note that only **write\_item** operations need to be undone during transaction rollback; **read\_item** operations are recorded in the log only to determine whether cascading rollback of additional transactions is necessary.

(b)

	A	B	C	D
	30	15	40	20
	[start_transaction, T <sub>3</sub> ]			
	[read_item, T <sub>3</sub> , C]			
*	[write_item, T <sub>3</sub> , B, 15, 12]			
	[start_transaction, T <sub>2</sub> ]			
	[read_item, T <sub>2</sub> , B]			
**	[write_item, T <sub>2</sub> , B, 12, 18]			
	[start_transaction, T <sub>1</sub> ]			
	[read_item, T <sub>1</sub> , A]			
	[read_item, T <sub>1</sub> , D]			
	[write_item, T <sub>1</sub> , D, 20, 25]			
	[read_item, T <sub>2</sub> , D]			
**	[write_item, T <sub>2</sub> , D, 25, 26]			
	[read_item, T <sub>3</sub> , A]			

← System crash

transaction (b) by the point of crash. (c) Operations before the crash.

\* T<sub>3</sub> is rolled back because it did not reach its commit point.

\*\* T<sub>2</sub> is rolled back because it reads the value of item B written by T<sub>3</sub>.