

```
import pandas as pd
import numpy as np
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, LSTM, Dense, SpatialDropout1D
from sklearn.model_selection import train_test_split
```

```
train_df = pd.read_csv('train.csv')
test_df = pd.read_csv('test.csv')
```

```
print(train_df.head())
print(test_df.head())
```

```
   id keyword location      text \
0    1    NaN     NaN  Our Deeds are the Reason of this #earthquake M...
1    4    NaN     NaN           Forest fire near La Ronge Sask. Canada
2    5    NaN     NaN  All residents asked to 'shelter in place' are ...
3    6    NaN     NaN  13,000 people receive #wildfires evacuation or...
4    7    NaN     NaN  Just got sent this photo from Ruby #Alaska as ...
```

```
   target
0        1
1        1
2        1
3        1
4        1
```

```
   id keyword location      text
0    0    NaN     NaN  Just happened a terrible car crash
1    2    NaN     NaN  Heard about #earthquake is different cities, s...
2    3    NaN     NaN  there is a forest fire at spot pond, geese are...
3    9    NaN     NaN  Apocalypse lighting. #Spokane #wildfires
4   11    NaN     NaN  Typhoon Soudelor kills 28 in China and Taiwan
```

```
tokenizer = Tokenizer()
tokenizer.fit_on_texts(train_df['text'])
```

```
X = tokenizer.texts_to_sequences(train_df['text'])
X = pad_sequences(X, maxlen=50)
```

```
y = train_df['target']
```

```
X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
model = Sequential()
model.add(Embedding(len(tokenizer.word_index) + 1, 32, input_length=50))
model.add(SpatialDropout1D(0.2))
model.add(LSTM(64, dropout=0.2, recurrent_dropout=0.2))
model.add(Dense(1, activation='sigmoid'))
```

```
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
```

```
print(model.summary())
```

```
Model: "sequential"
```

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 50, 32)	726432
spatial_dropout1d (Spatial Dropout1D)	(None, 50, 32)	0
lstm (LSTM)	(None, 64)	24832
dense (Dense)	(None, 1)	65

```
=====
Total params: 751329 (2.87 MB)
Trainable params: 751329 (2.87 MB)
Non-trainable params: 0 (0.00 Byte)
```

```
None
```

```
history = model.fit(X_train, y_train, epochs=10, batch_size=32, validation_data=(X_val, y_val), verbose=1)
```

```
Epoch 1/10
191/191 [=====] - 24s 94ms/step - loss: 0.5757 - accuracy: 0.6882 - val_loss: 0.4458 - val_accuracy: 0.812
```

```
Epoch 2/10
191/191 [=====] - 20s 103ms/step - loss: 0.3346 - accuracy: 0.8670 - val_loss: 0.4504 - val_accuracy: 0.80
Epoch 3/10
191/191 [=====] - 18s 92ms/step - loss: 0.2027 - accuracy: 0.9284 - val_loss: 0.5307 - val_accuracy: 0.799
Epoch 4/10
191/191 [=====] - 19s 97ms/step - loss: 0.1187 - accuracy: 0.9588 - val_loss: 0.6628 - val_accuracy: 0.774
Epoch 5/10
191/191 [=====] - 19s 100ms/step - loss: 0.0840 - accuracy: 0.9731 - val_loss: 0.8857 - val_accuracy: 0.74
Epoch 6/10
191/191 [=====] - 18s 97ms/step - loss: 0.0449 - accuracy: 0.9857 - val_loss: 0.8703 - val_accuracy: 0.741
Epoch 7/10
191/191 [=====] - 20s 102ms/step - loss: 0.0335 - accuracy: 0.9910 - val_loss: 0.8966 - val_accuracy: 0.75
Epoch 8/10
191/191 [=====] - 19s 97ms/step - loss: 0.0296 - accuracy: 0.9906 - val_loss: 0.8908 - val_accuracy: 0.747
Epoch 9/10
191/191 [=====] - 19s 99ms/step - loss: 0.0220 - accuracy: 0.9938 - val_loss: 0.9122 - val_accuracy: 0.753
Epoch 10/10
191/191 [=====] - 18s 95ms/step - loss: 0.0217 - accuracy: 0.9931 - val_loss: 0.8650 - val_accuracy: 0.726

loss, accuracy = model.evaluate(X_val, y_val)
print("Validation Accuracy:", accuracy)

48/48 [=====] - 1s 29ms/step - loss: 0.8650 - accuracy: 0.7262
Validation Accuracy: 0.7261983156204224

X_test = tokenizer.texts_to_sequences(test_df['text'])
X_test = pad_sequences(X_test, maxlen=50)

predictions_proba = model.predict(X_test).flatten()

102/102 [=====] - 1s 11ms/step

predictions = (predictions_proba > 0.5).astype(int)

submission_df = pd.DataFrame({'id': test_df['id'], 'target': predictions})
submission_df.to_csv('submission.csv', index=False)
```